



..

- sample_data
- hr_modified_1.csv

```
# upload file
# have removed spaces e.g "Reasearch & Development" to "RD"
import pandas as pd
import numpy as np
import seaborn as sns
from sklearn.impute import SimpleImputer
import matplotlib.pyplot as plt
from sklearn import preprocessing
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import OneHotEncoder
```

```
df_org = pd.read_csv("hr_modified_1.csv")
df = df_org
```

```
#replace space with nan
df = df.replace(r'^\s+$', np.nan, regex=True)
```

```
#df.info()
df[df.isna().any(axis=1)]
```

```

Age  Attrition  BusinessTravel  DailyRate  Department  DistanceFromHome
162   27         No           Travel_Rarely         NaN         RAD             2
516   33         No                NaN         1392         RAD             3
798   37         Yes                NaN         1373         RAD             2
1273  49         No   Travel_Frequently         279         RAD             8
```

4 rows × 35 columns



```
# df.info()
# listed are columns with nans
```

```
# BusinessTravel -char - 2
# Department - char - 1
# EducationField -char - 1
```

```
# DailyRate - numeric - 1
# MonthlyIncome -numeric - 1
# MonthlyRate -numeric - 1
```

```
# identifying BusinessTravel=='Travel_Rarely' and df.EducationField=='Medical' datas
df_medical_travel_rarely = df[(df.BusinessTravel=='Travel_Rarely') & (df.EducationFie
# replacing DailyRate with mean for dataset where BusinessTravel=='Travel_Rarely' and
df_daily_date = df_medical_travel_rarely['DailyRate']
df_df_daily_date = pd.DataFrame(df_daily_date)
df_df_daily_date['DailyRate'] = pd.to_numeric(df_df_daily_date['DailyRate'],errors =
```

```
# df_df_daily_date.mean()
df['DailyRate'] = df['DailyRate'].fillna(836.299694)
```

```
# df[df.isna().any(axis=1)]
#identifying Department for row 162 , Research And Development looks to be most occur
df_medical_travel_rarely.Department.value_counts(dropna=False)
```

```

RAD                255
Sales              62
Human_Resources   10
NaN                1
Name: Department, dtype: int64
```

```
df[df.isna().any(axis=1)]
```

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome
162	27	No	Travel_Rarely	836.299694	NaN	2
516	33	No	NaN	1392	RAD	3
798	37	Yes	NaN	1373	RAD	2
1273	49	No	Travel_Frequently	279	RAD	8

```
4 rows × 35 columns
# df[df.isna().any(axis=1)]
df['Department'] = df['Department'].fillna('RAD')

df.at[798, 'BusinessTravel'] = 'Travel_Rarely'

df.at[516, 'BusinessTravel'] = 'Travel_Rarely'


# checked for average of Attrition = Yes and Gender = Male and JobRole = Laboratory_T
df.at[798, 'MonthlyIncome'] = 2978

# checked for average of Attrition = No and Gender = Female and JobRole = Research_Sc
df.at[516, 'MonthlyRate'] = 15533

df
df.at[1273, 'EducationField'] = 'LS'

df[df.isna().any(axis=1)]
# df_no_missing_values = df
```

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Educa
0 rows × 35 columns							



```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 35 columns):
#   Column                                Non-Null Count  Dtype
---  ---                                -
0   Age                                  1470 non-null   int64
1   Attrition                          1470 non-null   object
2   BusinessTravel                     1470 non-null   object
3   DailyRate                          1470 non-null   object
4   Department                         1470 non-null   object
5   DistanceFromHome                   1470 non-null   int64
6   Education                          1470 non-null   int64
7   EducationField                     1470 non-null   object
8   EmployeeCount                      1470 non-null   int64
9   EmployeeNumber                     1470 non-null   int64
10  EnvironmentSatisfaction             1470 non-null   int64
11  Gender                             1470 non-null   object
12  HourlyRate                         1470 non-null   int64
13  JobInvolvement                     1470 non-null   int64
14  JobLevel                           1470 non-null   int64
15  JobRole                            1470 non-null   object
16  JobSatisfaction                     1470 non-null   int64
17  MaritalStatus                      1470 non-null   object
18  MonthlyIncome                      1470 non-null   object
19  MonthlyRate                        1470 non-null   object
20  NumCompaniesWorked                 1470 non-null   int64
21  Over18                             1470 non-null   object
22  OverTime                           1470 non-null   object
23  PercentSalaryHike                  1470 non-null   int64
24  PerformanceRating                  1470 non-null   int64
25  RelationshipSatisfaction             1470 non-null   int64
26  StandardHours                      1470 non-null   int64
27  StockOptionLevel                   1470 non-null   int64
28  TotalWorkingYears                  1470 non-null   int64
29  TrainingTimesLastYear              1470 non-null   int64
30  WorkLifeBalance                    1470 non-null   int64
31  YearsAtCompany                     1470 non-null   int64
32  YearsInCurrentRole                 1470 non-null   int64
33  YearsSinceLastPromotion             1470 non-null   int64
34  YearsWithCurrManager                1470 non-null   int64
dtypes: int64(23), object(12)
memory usage: 402.1+ KB
```

```
# columns listed below do not add any value on model.
# col = ['Attrition','EmployeeNumber','Over18']
col = ['EmployeeNumber','Over18']
X = df
for c in col:
    X = X.loc[:, X.columns != c]
y = df.Attrition
```

```
X['DailyRate'] = df['DailyRate'].astype(float)
X['MonthlyIncome'] = df['MonthlyIncome'].astype(float)
X['MonthlyRate'] = df['MonthlyRate'].astype(float)
```

```
X.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 33 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                   1470 non-null   int64
1   Attrition                           1470 non-null   object
2   BusinessTravel                       1470 non-null   object
3   DailyRate                           1470 non-null   float64
4   Department                           1470 non-null   object
5   DistanceFromHome                    1470 non-null   int64
6   Education                           1470 non-null   int64
7   EducationField                       1470 non-null   object
8   EmployeeCount                       1470 non-null   int64
9   EnvironmentSatisfaction              1470 non-null   int64
10  Gender                               1470 non-null   object
11  HourlyRate                           1470 non-null   int64
12  JobInvolvement                       1470 non-null   int64
13  JobLevel                             1470 non-null   int64
14  JobRole                              1470 non-null   object
15  JobSatisfaction                      1470 non-null   int64
16  MaritalStatus                       1470 non-null   object
17  MonthlyIncome                       1470 non-null   float64
18  MonthlyRate                         1470 non-null   float64
19  NumCompaniesWorked                  1470 non-null   int64
20  OverTime                            1470 non-null   object
21  PercentSalaryHike                   1470 non-null   int64
22  PerformanceRating                   1470 non-null   int64
23  RelationshipSatisfaction             1470 non-null   int64
24  StandardHours                       1470 non-null   int64
25  StockOptionLevel                    1470 non-null   int64
26  TotalWorkingYears                   1470 non-null   int64
27  TrainingTimesLastYear               1470 non-null   int64
28  WorkLifeBalance                     1470 non-null   int64
29  YearsAtCompany                      1470 non-null   int64
30  YearsInCurrentRole                  1470 non-null   int64
31  YearsSinceLastPromotion              1470 non-null   int64
32  YearsWithCurrManager                1470 non-null   int64
dtypes: float64(3), int64(22), object(8)
memory usage: 379.1+ KB
```

```
cat_cols=X.select_dtypes(include="object").columns
```

```
cat_cols_x = cat_cols.to_list()
cat_cols_x.remove('Attrition')
cat_cols_x
```

```
['BusinessTravel',
 'Department',
 'EducationField',
 'Gender',
 'JobRole',
 'MaritalStatus',
 'OverTime']
```

```
num_cols= X.select_dtypes(exclude="object").columns
```

```
num_cols
```

```
Index(['Age', 'DailyRate', 'DistanceFromHome', 'Education', 'EmployeeCount',
       'EnvironmentSatisfaction', 'HourlyRate', 'JobInvolvement', 'JobLevel',
       'JobSatisfaction', 'MonthlyIncome', 'MonthlyRate', 'NumCompaniesWorked',
       'PercentSalaryHike', 'PerformanceRating', 'RelationshipSatisfaction',
       'StandardHours', 'StockOptionLevel', 'TotalWorkingYears',
       'TrainingTimesLastYear', 'WorkLifeBalance', 'YearsAtCompany',
       'YearsInCurrentRole', 'YearsSinceLastPromotion',
```

```
'YearsWithCurrManager'],
dtype='object')
```


```
num_cols_x = num_cols.to_list()
no_need_to_scale = ['Education', 'EnvironmentSatisfaction', 'JobInvolvement', 'JobLevel',
for c in no_need_to_scale:
    num_cols_x.remove(c)
```

```
from sklearn.preprocessing import StandardScaler
sc= StandardScaler()
X[num_cols_x]= sc.fit_transform(X[num_cols_x])
```

```
y.replace('Yes',1, inplace=True)
y.replace('No',0, inplace=True)
```

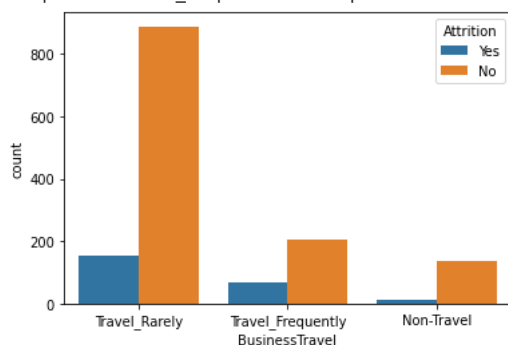
```
import seaborn as sns
sns.countplot (x=X.Department,hue=X.Attrition)
```

<matplotlib.axes._subplots.AxesSubplot at 0x7efda19551f0>



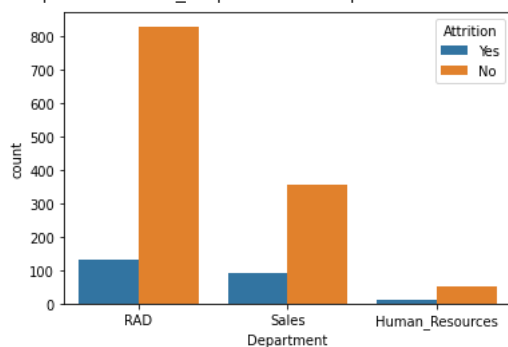
```
sns.countplot (x=X.BusinessTravel,hue=X.Attrition)
```

<matplotlib.axes._subplots.AxesSubplot at 0x7efd94ea5280>




```
sns.countplot (x=X.Department,hue=X.Attrition)
```

<matplotlib.axes._subplots.AxesSubplot at 0x7efd94e7ad60>




```
sns.countplot (x=X.EducationField,hue=X.Attrition)
```

<matplotlib.axes._subplots.AxesSubplot at 0x7efd94de79d0>



sns.countplot (x=X.Attrition)

<matplotlib.axes._subplots.AxesSubplot at 0x7efd94f95520>



Null accuracy = 1200/1400 = 85.7%

tried One Hot encoding with sklearn however the issue is i could not get "Column Name"
 # this causes issue for pairplots.

```
# from sklearn.preprocessing import OneHotEncoder
# from sklearn.compose import ColumnTransformer
```

```
# ct=ColumnTransformer(
# transformers=[('encoder',OneHotEncoder(sparse=False),[0]),remainder="passthrough")
# X_one=ct.fit_transform(X)
```

```
# pd.DataFrame(X_one)
# X_one
# ct.get_feature_names_out
```

```
# from sklearn.preprocessing import OneHotEncoder
# ohe = OneHotEncoder(sparse=False)
# X_ohe = ohe.fit_transform(X)
# ohe.get_feature_names_out()
```

```
X.drop(['Attrition'], axis=1, inplace=True)
dummies_df = pd.get_dummies(X, columns=cat_cols_x)
dummies_df.columns
dummies_df
```

	Age	DailyRate	DistanceFromHome	Education	EmployeeCount	Environm
0	-2.072192	-1.419794	-0.764121	3	0.0	
1	-2.072192	0.023175	0.099639	3	0.0	
2	-2.072192	1.247964	-0.517332	3	0.0	
3	-2.072192	-1.278472	-0.517332	2	0.0	
4	-2.072192	-1.377645	-0.147150	1	0.0	
...	
1465	2.526886	-0.943763	-0.270544	3	0.0	
1466	2.526886	1.726474	2.320735	3	0.0	
1467	2.526886	0.933089	0.840004	4	0.0	
1468	2.526886	-0.264427	-0.270544	4	0.0	
1469	2.526886	-1.072688	-1.010909	4	0.0	

1470 rows × 53 columns



```

from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(dummies_df,y,test_size=.25)

from sklearn.tree import DecisionTreeClassifier
clf = DecisionTreeClassifier()
clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)

dummies_df.columns

Index(['Age', 'DailyRate', 'DistanceFromHome', 'Education', 'EmployeeCount',
      'EnvironmentSatisfaction', 'HourlyRate', 'JobInvolvement', 'JobLevel',
      'JobSatisfaction', 'MonthlyIncome', 'MonthlyRate', 'NumCompaniesWorked',
      'PercentSalaryHike', 'PerformanceRating', 'RelationshipSatisfaction',
      'StandardHours', 'StockOptionLevel', 'TotalWorkingYears',
      'TrainingTimesLastYear', 'WorkLifeBalance', 'YearsAtCompany',
      'YearsInCurrentRole', 'YearsSinceLastPromotion', 'YearsWithCurrManager',
      'BusinessTravel_Non-Travel', 'BusinessTravel_Travel_Frequently',
      'BusinessTravel_Travel_Rarely', 'Department_Human_Resources',
      'Department_RAD', 'Department_Sales', 'EducationField_Human_Resources',
      'EducationField_LS', 'EducationField_Marketing',
      'EducationField_Medical', 'EducationField_Other',
      'EducationField_Technical_Degree', 'Gender_Female', 'Gender_Male',
      'JobRole_Healthcare_Representative', 'JobRole_Human_Resources',
      'JobRole_Laboratory_Technician', 'JobRole_Manager',
      'JobRole_Manufacturing_Director', 'JobRole_Research_Director',
      'JobRole_Research_Scientist', 'JobRole_Sales_Executive',
      'JobRole_Sales_Representative', 'MaritalStatus_Divorced',
      'MaritalStatus_Married', 'MaritalStatus_Single', 'OverTime_No',
      'OverTime_Yes'],
      dtype='object')

from sklearn.metrics import accuracy_score
print(accuracy_score(y_test, y_pred))

0.7336956521739131

from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test,y_pred)
cm

array([[252, 64],
       [ 34, 18]])

# !pip install sklearn.externals.six
# from sklearn.externals.six import StringIO
# from IPython.display import Image
# from sklearn.tree import export_graphviz
# import pydotplus

# dot_data = StringIO()
# export_graphviz(clf, out_file=dot_data,
#                 filled=True, rounded=True,
#                 special_characters=True)

# graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
# Image(graph.create_png())

# https://machinelearningmastery.com/calculate-feature-importance-with-python/
from matplotlib import pyplot
# get importance
importance = clf.feature_importances_
# summarize feature importance
for i,v in enumerate(importance):
    print('Feature: %0d, Score: %.5f' % (i,v))
# plot feature importance
pyplot.bar([x for x in range(len(importance))], importance)
pyplot.show()

```

Feature: 0, Score: 0.06610
Feature: 1, Score: 0.09201
Feature: 2, Score: 0.06557
Feature: 3, Score: 0.02035
Feature: 4, Score: 0.00000
Feature: 5, Score: 0.03107
Feature: 6, Score: 0.07556
Feature: 7, Score: 0.00879
Feature: 8, Score: 0.00561
Feature: 9, Score: 0.01515
Feature: 10, Score: 0.11230
Feature: 11, Score: 0.03029
Feature: 12, Score: 0.01608
Feature: 13, Score: 0.05648
Feature: 14, Score: 0.00433
Feature: 15, Score: 0.02221
Feature: 16, Score: 0.00000
Feature: 17, Score: 0.05441
Feature: 18, Score: 0.00325
Feature: 19, Score: 0.00709
Feature: 20, Score: 0.02838
Feature: 21, Score: 0.03732
Feature: 22, Score: 0.02209
Feature: 23, Score: 0.03477
Feature: 24, Score: 0.00000
Feature: 25, Score: 0.00000
Feature: 26, Score: 0.01895
Feature: 27, Score: 0.00000
Feature: 28, Score: 0.00839
Feature: 29, Score: 0.00000
Feature: 30, Score: 0.00000
Feature: 31, Score: 0.00000
Feature: 32, Score: 0.00000
Feature: 33, Score: 0.00000
Feature: 34, Score: 0.00780
Feature: 35, Score: 0.00000
Feature: 36, Score: 0.01482
Feature: 37, Score: 0.00487
Feature: 38, Score: 0.00000
Feature: 39, Score: 0.00000
Feature: 40, Score: 0.00000
Feature: 41, Score: 0.00711
Feature: 42, Score: 0.00000
Feature: 43, Score: 0.00000
Feature: 44, Score: 0.00000
Feature: 45, Score: 0.00585
Feature: 46, Score: 0.03017
Feature: 47, Score: 0.00000
Feature: 48, Score: 0.00000
Feature: 49, Score: 0.00000
Feature: 50, Score: 0.00000
Feature: 51, Score: 0.05918
Feature: 52, Score: 0.03365



Following come out to be important features from CART

1	0.07703 'DailyRate',
18	0.06661 'TotalWorkingYears',
10	0.06341 'MonthlyIncome',
21	0.0603 'YearsAtCompany',
8	0.06009 'JobLevel',
2	0.05733 'DistanceFromHome',
11	0.05533 'MonthlyRate',
0	0.04628 'Age',
9	0.04415 'JobSatisfaction',
51	0.04333 'OverTime_No',
12	0.04019 'NumCompaniesWorked',
13	0.03464 'PercentSalaryHike',
6	0.03378 'HourlyRate',
19	0.02847 'TrainingTimesLastYear',
20	0.02695 'WorkLifeBalance',
37	0.02613 'Gender_Female',
22	0.02596 'YearsInCurrentRole',
7	0.02153 'JobInvolvement',
52	0.02075 'OverTime_Yes'
23	0.02 'YearsSinceLastPromotion',

```
cat_cols_feat = dummies_df.select_dtypes(include="object").columns
num_cols_feat = dummies_df.select_dtypes(exclude="object").columns

from sklearn.pipeline import Pipeline
from sklearn.preprocessing import OneHotEncoder
from sklearn.linear_model import LogisticRegression
```

```

from sklearn.compose import ColumnTransformer

# Cat Tranformer
categorical_tranformer= Pipeline(steps=[('ohe',OneHotEncoder())])

# Num tranfoemr
numerical_tranformer = Pipeline(steps=[('sc',StandardScaler())])

col_tranform= ColumnTransformer(transformers=[
                                ('cat_feat',categorical_tranformer,cat_cols_fea
                                ('num_feat',numerical_tranformer,num_cols_feat),
                                ],
                                remainder='passthrough')

X_trans = col_tranform.fit_transform(dummies_df)
type(X_trans)

numpy.ndarray

my_pipeline= Pipeline(steps=[('first_pipe',col_tranform),('model',LogisticRegression(
# X_train

my_pipeline.fit(X_train,y_train)

Pipeline(steps=[('first_pipe',
                ColumnTransformer(remainder='passthrough',
                                transformers=[('cat_feat',
                                                Pipeline(steps=[('ohe',
                                OneHotEncoder()))]),
                                Index([], dtype='object')),
                                ('num_feat',
                                Pipeline(steps=[('sc',
                                StandardScaler()))]),
                                Index(['Age', 'DailyRate',
'DistanceFromHome', 'Education', 'EmployeeCount',
'EnvironmentSatisfaction', 'HourlyRate', 'JobInvolvement', 'J...
'JobRole_Laboratory_Technician', 'JobRole_Manager',
'JobRole_Manufacturing_Director', 'JobRole_Research_Director',
'JobRole_Research_Scientist', 'JobRole_Sales_Executive',
'JobRole_Sales_Representative', 'MaritalStatus_Divorced',
'MaritalStatus_Married', 'MaritalStatus_Single', 'OverTime_No',
'OverTime_Yes'],
dtype='object')))])),
                ('model', LogisticRegression()))])

y_pred= my_pipeline.predict(X_test)

from sklearn.metrics import accuracy_score
pd.Series(accuracy_score(y_test,y_pred))

0    0.88587
dtype: float64

from sklearn import metrics
metrics.confusion_matrix(y_test,y_pred)

array([[301, 15],
       [ 27, 25]])

metrics.recall_score(y_test,y_pred)

0.4807692307692308

from sklearn.ensemble import RandomForestClassifier
my_pipeline= Pipeline(steps=[('first_pipe',col_tranform),('model',RandomForestClassif
my_pipeline.fit(X_train,y_train)

Pipeline(steps=[('first_pipe',
                ColumnTransformer(remainder='passthrough',
                                transformers=[('cat_feat',
                                                Pipeline(steps=[('ohe',
                                OneHotEncoder()))]),
                                Index([], dtype='object')),
                                ('num_feat',
                                Pipeline(steps=[('sc',
                                StandardScaler()))]),
                                Index(['Age', 'DailyRate',

```



```

'DistanceFromHome', 'Education', 'EmployeeCount',
  'EnvironmentSatisfaction', 'HourlyRate', 'JobInvolvement', 'J...
  'JobRole_Laboratory_Technician', 'JobRole_Manager',
  'JobRole_Manufacturing_Director', 'JobRole_Research_Director',
  'JobRole_Research_Scientist', 'JobRole_Sales_Executive',
  'JobRole_Sales_Representative', 'MaritalStatus_Divorced',
  'MaritalStatus_Married', 'MaritalStatus_Single', 'OverTime_No',
  'OverTime_Yes'],
dtype='object')))),
    ('model', RandomForestClassifier()))

my_pipeline.fit(X_train,y_train)
y_pred= my_pipeline.predict(X_test)

from sklearn.metrics import accuracy_score
pd.Series(accuracy_score(y_test,y_pred))

0    0.877717
dtype: float64

from sklearn import metrics
metrics.confusion_matrix(y_test,y_pred)

array([[313,   3],
       [ 42,  10]])

metrics.precision_score(y_test,y_pred)

0.7692307692307693

metrics.recall_score(y_test,y_pred)

0.19230769230769232

metrics.f1_score(y_test,y_pred)

0.3076923076923077

from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import RandomForestClassifier

# # Define Parameters
# max_depth=[2, 8, 16]
# n_estimators = [64, 128, 256]
# param_grid = dict(max_depth=max_depth, n_estimators=n_estimators)

# # Build the grid search
# dfrst = RandomForestClassifier(n_estimators=n_estimators, max_depth=max_depth)
# grid = GridSearchCV(estimator=dfrst, param_grid=param_grid, cv = 5, scoring = 'accuracy')
# grid_results = grid.fit(X_train, y_train)

# # Summarize the results in a readable format
# print("Best: {0}, using {1}".format(grid_results.cv_results_['accuracy'], grid_results.best_index_))
# results_df = pd.DataFrame(grid_results.cv_results_)
# results_df

# rf = RandomForestClassifier()

# # grid search cv
# grid_space={'max_depth':[3,5,10,None],
#             'n_estimators':[10,100,200],
#             'max_features':[1,3,5,7],
#             'min_samples_leaf':[1,2,3],
#             'min_samples_split':[1,2,3]
#             }

# grid = GridSearchCV(rf,param_grid=grid_space,cv=3,scoring='accuracy')
# model_grid = grid.fit(X_train,y_train)

# # grid search results
# print('Best grid search hyperparameters are: '+str(model_grid.best_params_))
# print('Best grid search score is: '+str(model_grid.best_score_))

!pip install sweetviz
import sweetviz

```

```

Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wh
Requirement already satisfied: sweetviz in /usr/local/lib/python3.8/dist-packag
Requirement already satisfied: importlib-resources>=1.2.0 in /usr/local/lib/pyt
Requirement already satisfied: matplotlib>=3.1.3 in /usr/local/lib/python3.8/di
Requirement already satisfied: Jinja2>=2.11.1 in /usr/local/lib/python3.8/dist-
Requirement already satisfied: scipy>=1.3.2 in /usr/local/lib/python3.8/dist-pa
Requirement already satisfied: tqdm>=4.43.0 in /usr/local/lib/python3.8/dist-pa
Requirement already satisfied: pandas!=1.0.0,!=1.0.1,!=1.0.2,>=0.25.3 in /usr/l
Requirement already satisfied: numpy>=1.16.0 in /usr/local/lib/python3.8/dist-p
Requirement already satisfied: zipp>=3.1.0 in /usr/local/lib/python3.8/dist-pac
Requirement already satisfied: MarkupSafe>=0.23 in /usr/local/lib/python3.8/dis
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /usr
Requirement already satisfied: cyclers>=0.10 in /usr/local/lib/python3.8/dist-pa
Requirement already satisfied: python-dateutil>=2.1 in /usr/local/lib/python3.8
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.8/di
Requirement already satisfied: pytz>=2017.3 in /usr/local/lib/python3.8/dist-pa
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.8/dist-packag

```

```

# my_report = sweetviz.compare([X, "Train"] "Attrition")
# my_report =sweetviz.analyze(source=df_org,target_feat='Attrition')

```

```
# my_report.show_html("Report.html")
```

**** Gaussian naive bayes ** this is the best fit of all**

```

# X_train,X_test,y_train,y_test
from sklearn.naive_bayes import GaussianNB
GNBclf = GaussianNB()
model = GNBclf.fit(X_train, y_train)
y_pred = model.predict(X_test)

```

```
metrics.recall_score(y_test,y_pred)
```

```
0.7115384615384616
```

```
metrics.precision_score(y_test,y_pred)
```

```
0.25170068027210885
```

```
metrics.f1_score(y_test,y_pred)
```

```
0.37185929648241206
```

```
metrics.confusion_matrix(y_test,y_pred)
```

```
array([[206, 110],
       [ 15,  37]])
```

```


from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
metrics.confusion_matrix(y_test,y_pred)
# metrics.recall_score(y_test,y_pred)
cm = confusion_matrix(y_test, y_pred, labels=clf.classes_)
disp = ConfusionMatrixDisplay(confusion_matrix=cm,
                              display_labels=clf.classes_)
disp.plot()

```

```

<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at
0x7efd9c4b1e50>

```




```

from numpy import argmax
from sklearn.metrics import plot_roc_curve
roc = plot_roc_curve(model, X_test, y_test)

```

/usr/local/lib/python3.8/dist-packages/sklearn/utils/deprecation.py:87: FutureWarning: warnings.warn(msg, category=FutureWarning)



```
y_pred_prob = model.predict_proba(X_test)[: ,1]
```

```
y_pred_prob
```

```
2.92002670e-05, 2.76470963e-03, 1.31719269e-17, 6.73637676e-01,
9.96560115e-01, 8.29874625e-13, 8.70571987e-01, 9.96415717e-01,
6.96922333e-02, 8.03420852e-01, 1.62506069e-04, 4.95565050e-01,
2.87049549e-04, 8.78774642e-05, 1.27389193e-01, 3.05996456e-02,
7.92601619e-01, 1.55333892e-01, 3.60842166e-04, 9.88196492e-01,
9.81841915e-04, 2.34407597e-45, 1.11041725e-04, 1.24308234e-01,
7.38001888e-07, 3.00622466e-14, 4.68746585e-01, 9.69055437e-01,
9.99703617e-01, 9.24536918e-01, 4.59105778e-08, 4.52572990e-39,
2.36668103e-06, 9.97039016e-01, 9.29786616e-03, 9.31145361e-10,
2.09595470e-43, 8.62356416e-01, 1.64683281e-07, 1.10359321e-40,
7.58234011e-04, 9.38404343e-01, 7.67031960e-01, 9.99998784e-01,
9.61046482e-01, 1.10745361e-42, 1.42515397e-02, 3.54653565e-03,
8.28224104e-01, 4.63347459e-19, 7.61644690e-01, 8.57443457e-01,
9.92789120e-01, 9.9997591e-01, 6.24181506e-02, 9.99974950e-01,
9.99999610e-01, 4.65120980e-01, 2.06160444e-12, 5.35875698e-01,
2.15661537e-02, 6.54087677e-03, 1.63745887e-02, 3.63879109e-10,
4.01967434e-04, 3.98726152e-08, 1.01332510e-01, 3.32597937e-41,
9.95906340e-01, 6.17867233e-01, 4.76788213e-01, 1.57245312e-16,
1.39261951e-39, 6.53084330e-08, 3.66512600e-02, 9.99999962e-01,
5.47572038e-01, 3.14351175e-07, 1.10615369e-01, 9.97903631e-01,
2.30076083e-01, 1.47723295e-01, 9.77505617e-01, 9.99345438e-01,
9.83873092e-01, 3.42122825e-07, 1.38338784e-43, 9.99745825e-01,
1.36977564e-03, 3.70122251e-16, 1.40451114e-41, 4.78597450e-06,
6.85867918e-04, 9.13594565e-02, 1.19694528e-01, 2.29196568e-42,
2.78150005e-02, 1.72588120e-04, 9.82773558e-01, 9.65937060e-01,
9.91786794e-01, 1.80912073e-06, 8.26930003e-01, 1.98495155e-15,
1.80894966e-01, 9.43889204e-01, 9.39557314e-02, 6.64074663e-06,
9.90204099e-01, 6.54753124e-01, 4.24314004e-02, 9.07241235e-01,
9.58881698e-07, 1.61409296e-03, 2.55808527e-19, 2.93798674e-05,
5.85966764e-01, 1.60004222e-04, 9.98478941e-01, 1.51521245e-02,
7.15443293e-03, 7.32402621e-07, 1.87017944e-03, 8.70875696e-01,
1.68789278e-01, 9.76159501e-01, 2.02915232e-40, 2.83695382e-07,
9.48667882e-01, 1.40773580e-19, 1.07823205e-03, 8.81682064e-01,
4.46441865e-44, 2.04105354e-05, 9.60613367e-01, 6.11504540e-04,
2.02543374e-02, 1.16420107e-01, 1.67542057e-03, 3.59313004e-01,
3.32365970e-01, 6.92834219e-04, 9.98087855e-01, 1.76200596e-01,
6.43208447e-02, 3.42387211e-06, 2.45052921e-01, 3.02991933e-05,
8.05800237e-01, 3.82873051e-16, 9.46490825e-01, 6.66973730e-01,
6.25620561e-07, 4.43688577e-01, 4.27236924e-44, 6.75321779e-20,
6.73530241e-01, 9.99418741e-01, 7.64761553e-01, 6.97478310e-01,
2.90798735e-03, 4.02793650e-42, 4.55574016e-10, 4.72018588e-01,
1.90632459e-01, 8.76387166e-01, 7.63871285e-01, 2.68990462e-01,
9.63846200e-01, 6.40349178e-01, 9.87829345e-01, 1.11274476e-03,
9.98560038e-01, 3.20039453e-12, 3.62807567e-02, 4.58669892e-01,
1.09112950e-01, 1.58309372e-07, 5.04427204e-07, 9.63841756e-01,
6.82147313e-01, 5.99393793e-17, 3.01688974e-01, 9.99999995e-01,
6.63483724e-01, 6.63301734e-13, 1.52142015e-06, 1.39567167e-01,
8.43503483e-01, 5.12323044e-01, 9.99901402e-01, 7.54094866e-01,
6.47545256e-07, 2.15976502e-01, 3.04738419e-01, 9.89597685e-01,
9.53320231e-01, 1.11113802e-40, 4.51665067e-01, 9.99998614e-01,
8.69879335e-01, 9.99442139e-01, 1.12021546e-02, 1.53756696e-01,
6.26364445e-01, 5.20946366e-01, 5.92528415e-02, 5.63400550e-07,
8.09149272e-01, 9.97225615e-01, 9.93946040e-01, 8.62359644e-02,
6.53541422e-01, 2.94899158e-01, 9.76651182e-01, 8.99417544e-01,
1.14268322e-06, 7.64728586e-08, 1.85814569e-01, 3.21990546e-07,
2.66540017e-02, 8.96057027e-01, 2.62364618e-01, 6.65928683e-01,
6.28303843e-45, 5.70863676e-01, 1.23020917e-03, 5.36148760e-01,
6.82893115e-03, 6.99930787e-07, 8.37429588e-01, 7.61274252e-20]]
```

```
dummies_df.columns
```


```
Index(['Age', 'DailyRate', 'DistanceFromHome', 'Education', 'EmployeeCount',
       'EnvironmentSatisfaction', 'HourlyRate', 'JobInvolvement', 'JobLevel',
       'JobSatisfaction', 'MonthlyIncome', 'MonthlyRate', 'NumCompaniesWorked',
       'PercentSalaryHike', 'PerformanceRating', 'RelationshipSatisfaction',
```

```
'StandardHours', 'StockOptionLevel', 'TotalWorkingYears',
'TrainingTimesLastYear', 'WorkLifeBalance', 'YearsAtCompany',
'YearsInCurrentRole', 'YearsSinceLastPromotion', 'YearsWithCurrManager',
'BusinessTravel_Non-Travel', 'BusinessTravel_Travel_Frequently',
'BusinessTravel_Travel_Rarely', 'Department_Human_Resources',
'Department_RAD', 'Department_Sales', 'EducationField_Human_Resources',
'EducationField_LS', 'EducationField_Marketing',
'EducationField_Medical', 'EducationField_Other',
'EducationField_Technical_Degree', 'Gender_Female', 'Gender_Male',
'JobRole_Healthcare_Representative', 'JobRole_Human_Resources',
'JobRole_Laboratory_Technician', 'JobRole_Manager',
'JobRole_Manufacturing_Director', 'JobRole_Research_Director',
'JobRole_Research_Scientist', 'JobRole_Sales_Executive',
'JobRole_Sales_Representative', 'MaritalStatus_Divorced',
'MaritalStatus_Married', 'MaritalStatus_Single', 'OverTime_No',
'OverTime_Yes'],
dtype='object')
```

Trying to identify if 'DailyRate','HourlyRate','MonthlyRate','MonthlyIncome' are co-related and any of these can be eliminated. Checking this since Naive bayes treats each dimension independently and would work better if they are not related.

```
dummies_df_corr = dummies_df[['DailyRate','HourlyRate','MonthlyRate','MonthlyIncome']]
```

```
correlations = dummies_df_corr.corr()
#sns.heatmap(correlations, xticklabels=correlations.columns, yticklabels=correlations.columns)
sns.clustermap(correlations, xticklabels=correlations.columns, yticklabels=correlations.columns)
```



Do not find strong correlation between these columns - 'DailyRate','HourlyRate','MonthlyRate','MonthlyIncome'. no Pearson's correlation value is above 0.5 Trying to eliminate correlated columns for Naive bayes.

```
dummies_df.info()
```

Data columns (total 53 columns):			
#	Column	Non-Null Count	Dtype
---	-----	-----	-----

```

2 DistanceFromHome      1470 non-null float64
3 Education              1470 non-null int64
4 EmployeeCount          1470 non-null float64
5 EnvironmentSatisfaction 1470 non-null int64
6 HourlyRate             1470 non-null float64
7 JobInvolvement         1470 non-null int64
8 JobLevel               1470 non-null int64
9 JobSatisfaction        1470 non-null int64
10 MonthlyIncome         1470 non-null float64
11 MonthlyRate           1470 non-null float64
12 NumCompaniesWorked    1470 non-null float64
13 PercentSalaryHike     1470 non-null int64
14 PerformanceRating     1470 non-null int64
15 RelationshipSatisfaction 1470 non-null int64
16 StandardHours         1470 non-null float64
17 StockOptionLevel      1470 non-null int64
18 TotalWorkingYears     1470 non-null float64
19 TrainingTimesLastYear 1470 non-null float64
20 WorkLifeBalance       1470 non-null int64
21 YearsAtCompany        1470 non-null float64
22 YearsInCurrentRole    1470 non-null float64
23 YearsSinceLastPromotion 1470 non-null float64
24 YearsWithCurrManager  1470 non-null float64
25 BusinessTravel_Non-Travel 1470 non-null uint8
26 BusinessTravel_Travel_Frequently 1470 non-null uint8
27 BusinessTravel_Travel_Rarely 1470 non-null uint8
28 Department_Human_Resources 1470 non-null uint8
29 Department_RAD        1470 non-null uint8
30 Department_Sales      1470 non-null uint8
31 EducationField_Human_Resources 1470 non-null uint8
32 EducationField_LS     1470 non-null uint8
33 EducationField_Marketing 1470 non-null uint8
34 EducationField_Medical 1470 non-null uint8
35 EducationField_Other   1470 non-null uint8
36 EducationField_Technical_Degree 1470 non-null uint8
37 Gender_Female         1470 non-null uint8
38 Gender_Male           1470 non-null uint8
39 JobRole_Healthcare_Representative 1470 non-null uint8
40 JobRole_Human_Resources 1470 non-null uint8
41 JobRole_Laboratory_Technician 1470 non-null uint8
42 JobRole_Manager       1470 non-null uint8
43 JobRole_Manufacturing_Director 1470 non-null uint8
44 JobRole_Research_Director 1470 non-null uint8
45 JobRole_Research_Scientist 1470 non-null uint8
46 JobRole_Sales_Executive 1470 non-null uint8
47 JobRole_Sales_Representative 1470 non-null uint8
48 MaritalStatus_Divorced 1470 non-null uint8
49 MaritalStatus_Married 1470 non-null uint8
50 MaritalStatus_Single  1470 non-null uint8
51 OverTime_No           1470 non-null uint8
52 OverTime_Yes          1470 non-null uint8
dtypes: float64(15), int64(10), uint8(28)
memory usage: 327.4 KB

```

```
dummies_df_corr = dummies_df[['DailyRate', 'HourlyRate', 'MonthlyRate', 'MonthlyIncome']]
```

```

from scipy.stats import pearsonr
# calculate Pearson's correlation
# corr, _ = pearsonr(dummies_df_corr['DailyRate'], dummies_df_corr['HourlyRate'])
corr, _ = pearsonr(dummies_df_corr['DailyRate'], dummies_df_corr['HourlyRate'])
print('Pearsons correlation: %.3f' % corr)

```

```
Pearsons correlation: 0.023
```

```

param_grid_nb = {
    'var_smoothing': np.logspace(0, -9, num=100)
}

```

```

from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import GridSearchCV
nbModel_grid = GridSearchCV(estimator=GaussianNB(), param_grid=param_grid_nb, verbose=1)
nbModel_grid.fit(X_test, y_test)
print(nbModel_grid.best_estimator_)

```

```

Fitting 10 folds for each of 100 candidates, totalling 1000 fits
GaussianNB(var_smoothing=0.0657933224657568)

```

```
y_pred = nbModel_grid.predict(X_test)
```

```

from sklearn.metrics import confusion_matrix
print(confusion_matrix(y_test, y_pred), ": is the confusion matrix")

```

```

from sklearn.metrics import f1_score
print(f1_score(y_test, y_pred), ": is the f1 score")


```

```
from sklearn.metrics import recall_score
print(recall_score(y_test, y_pred), ": is the recall score")
```

```
[[316   0]
 [ 46   6]] : is the confusion matrix
0.20689655172413793 : is the f1 score
0.11538461538461539 : is the recall score
```

```
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
metrics.confusion_matrix(y_test,y_pred)
# metrics.recall_score(y_test,y_pred)
cm = confusion_matrix(y_test, y_pred, labels=clf.classes_)
disp = ConfusionMatrixDisplay(confusion_matrix=cm,
                              display_labels=clf.classes_)
disp.plot()
```

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at
0x7efd94bc2c40>
```




```
from sklearn.metrics import f1_score
print(f1_score(y_test, y_pred), ": is the f1 score")
```

```
0.20689655172413793 : is the f1 score
```

```
roc = plot_roc_curve(model, X_test, y_pred)
```

```
/usr/local/lib/python3.8/dist-packages/sklearn/utils/deprecation.py:87: FutureWarning:
warnings.warn(msg, category=FutureWarning)
```




```
from numpy import mean
from sklearn.datasets import make_classification
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import cross_val_predict
from sklearn.model_selection import RepeatedStratifiedKFold
from imblearn.pipeline import Pipeline
from imblearn.over_sampling import SMOTE

# define pipeline
steps = [('over', SMOTE()), ('model', GaussianNB())]
pipeline = Pipeline(steps=steps)
# evaluate pipeline
cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
# scores = cross_val_score(pipeline, X_train, y_train, scoring='roc_auc', cv=cv, n_jobs=4)
# print('Mean ROC AUC: %.3f' % mean(scores))
# scores = cross_val_score(pipeline, X_train, y_train, scoring='recall', cv=cv, n_jobs=4)
# print('Mean Recall: %.3f' % mean(scores))
my_pipeline= Pipeline(steps=[('over', SMOTE()), ('model', GaussianNB())])
my_pipeline.fit(X_train,y_train)
y_pred= my_pipeline.predict(X_test)
from sklearn.metrics import accuracy_score
# pd.Series(accuracy_score(y_test,y_pred))
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
metrics.confusion_matrix(y_test,y_pred)
```

```
# metrics.recall_score(y_test,y_pred)
cm = confusion_matrix(y_test, y_pred, labels=clf.classes_)
disp = ConfusionMatrixDisplay(confusion_matrix=cm,
                              display_labels=clf.classes_)
disp.plot()
```

<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x7efd9513e790>



```
from sklearn.metrics import f1_score
print(f1_score(y_test, y_pred), ": is the f1 score")
```

0.3298969072164949 : is the f1 score

```
from sklearn.metrics import recall_score
print(recall_score(y_test, y_pred), ": is the recall score")
```

0.6153846153846154 : is the recall score

Double-click (or enter) to edit

Applying PCA - this is expected to give better results for Naive Bayes

```
# Set the n_components=3
from sklearn.decomposition import PCA


principal=PCA(n_components=18)
principal.fit(dummies_df)
Principal_components =principal.transform(dummies_df)

df_pca = pd.DataFrame(data = Principal_components, columns = ['PC 1', 'PC 2', 'PC 3',
# Check the dimensions of data after PCA
print(df_pca.shape)

# check how much variance is explained by each principal component
print(principal.explained_variance_ratio_)

plt.plot(principal.explained_variance_ratio_.cumsum(), marker='p', color='r', ls='-')
```

```
(1470, 18)
[0.35372911 0.12762878 0.04834421 0.03329 0.03232717 0.03083077
0.0293474 0.02798081 0.02747513 0.02583569 0.02461236 0.02439993
0.02255323 0.02089638 0.01732752 0.01426953 0.01393021 0.01299408]
[<matplotlib.lines.Line2D at 0x7efd945f77f0>]
```



df_pca


```

0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0])

```

```

from sklearn.metrics import confusion_matrix
print(confusion_matrix(y_test, y_pred), ": is the confusion matrix")

```

```

from sklearn.metrics import f1_score
print(f1_score(y_test, y_pred), ": is the f1 score")

```

```

from sklearn.metrics import recall_score
print(recall_score(y_test, y_pred), ": is the recall score")

```

```

[[305  3]
 [ 51  9]] : is the confusion matrix
0.24999999999999997 : is the f1 score
0.15 : is the recall score

```

```

from numpy import mean
from sklearn.datasets import make_classification
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import cross_val_predict
from sklearn.model_selection import RepeatedStratifiedKFold
from imblearn.pipeline import Pipeline
from imblearn.over_sampling import SMOTE


# define pipeline
steps = [('over', SMOTE()), ('model', GaussianNB())]
pipeline = Pipeline(steps=steps)
# evaluate pipeline
cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
# scores = cross_val_score(pipeline, X_train, y_train, scoring='roc_auc', cv=cv, n_jobs=
# print('Mean ROC AUC: %.3f' % mean(scores))
# scores = cross_val_score(pipeline, X_train, y_train, scoring='recall', cv=cv, n_jobs=
# print('Mean Recall: %.3f' % mean(scores))
my_pipeline= Pipeline(steps=[('over', SMOTE()), ('model', GaussianNB())])
my_pipeline.fit(X_train,y_train)
y_pred= my_pipeline.predict(X_test)
from sklearn.metrics import accuracy_score
# pd.Series(accuracy_score(y_test,y_pred))
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
metrics.confusion_matrix(y_test,y_pred)
# metrics.recall_score(y_test,y_pred)
cm = confusion_matrix(y_test, y_pred, labels=clf.classes_)
disp = ConfusionMatrixDisplay(confusion_matrix=cm,
                             display_labels=clf.classes_)
disp.plot()

```

```

<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at
0x7efd945248b0>

```



```

from sklearn.metrics import f1_score
print(f1_score(y_test, y_pred), ": is the f1 score")

```

```

0.3567567567567568 : is the f1 score

```

Disk 83.31 GB available

✓ 0s completed at 12:47 PM

